

AI VIDEO BLUEPRINT

How to Make Cinematic PS2-Style AI Videos From Scratch

A practical, no-fluff workflow for turning an idea into a finished short-form video — using AI tools to handle the execution while you handle the creativity.

00 — START HERE

What this Blueprint teaches

This Blueprint teaches you the complete process for making short, funny, cinematic AI-generated videos in a PS2-era, early-2000s Japanese video-game visual style — from a blank idea to a finished, edited video you can post.

By the end, you'll know how to:

- Design a character using AI image generation
- Build a small library of reference images so that character stays consistent
- Create a matching environment/background
- Write a short, punchy script yourself (not outsource it to AI)
- Plan shots like a director, even with zero filmmaking background
- Turn still images into short video clips using AI
- Prompt that AI with enough specificity that it actually does what you want
- Fix a generation when it goes wrong instead of starting over
- Generate and direct a voice performance
- Edit everything together into a finished video in CapCut

Who this is for

This is written for a complete beginner. No filmmaking background, no 3D modeling experience, and no prior AI tool experience is assumed. If you can type a sentence and follow numbered steps, you can follow this Blueprint.

What you'll need

- An AI image-generation tool (this Blueprint uses **ChatGPT image generation** as the default recommendation, since it's simple and accessible — any comparable image generator will work)

- An AI image-to-video tool (this Blueprint is built around **Google Flow**, using whatever Veo/Flow model is currently available to you)
- A video editor (this Blueprint uses **CapCut**)
- Optionally, a dedicated voice-generation tool such as **ElevenLabs**, if you want more control over voice than your video tool's built-in voice offers

You do not need a expensive computer, a camera, or any physical filming equipment. Everything here is generated.

The core philosophy

You come up with the idea. AI handles the execution. Your job is to tell the AI exactly what you want.

This idea runs through the entire Blueprint. AI tools are extremely good at execution once you give them clear direction. They are not good at being funny, being original, or reading your mind. That part is yours.

CREATOR TIP: The people who get the best results from these tools aren't the ones with the fanciest tools. They're the ones who are specific. Specificity is the actual skill you're learning here — everything else is just software.

The overall workflow

IDEA → CHARACTER → REFERENCE IMAGES → ENVIRONMENT → SCRIPT → SHOT PLANNING
 → IMAGE-TO-VIDEO → VOICE/PERFORMANCE → MULTIPLE SHORT CLIPS
 → CAPCUT EDITING → FINAL VIDEO

Every chapter in this Blueprint maps to one step in that chain. Read them in order the first time. After that, use it as a reference and jump to whatever step you're stuck on.

The role of AI vs. human creativity

A simple way to think about it:

You do	AI does
The idea, the joke, the premise	Generating the character image
The script	Generating the environment
The shot decisions	Turning images into motion
Deciding what's wrong with a generation	Generating a voice

Final creative judgment

Doing exactly what you tell it to

If you remember nothing else from this Blueprint, remember this: **a technically perfect AI video with a boring idea is still a boring video**. Don't skip Chapter 01.

01 — THE IDEA

Why the idea matters more than the technology

It is tempting to think the hard part of this process is the AI tools. It isn't. Once you understand the workflow, generating a character, an environment, and a clip is mechanical. The hard part — the part that actually determines whether anyone watches your video — is the idea.

A technically flawless AI video with a boring premise will still get scrolled past. A rough-looking video with a genuinely funny idea will still land.

This is good news. It means you don't need the newest AI model or the most polished visuals to succeed. You need a decent idea and a workflow to execute it. This Blueprint gives you the workflow. The idea is on you.

Finding a premise

You don't need a complicated concept. Short-form comedy usually works best when it's simple enough to explain in one sentence. A useful test: if you can't describe your idea in one sentence to a friend and get a reaction, it's probably not ready yet.

Some directions that tend to work well for this style of video:

- A character reacting to something mundane as if it's a huge deal (mock-seriousness)
- A character stating something absurd with complete sincerity
- A familiar situation (a job interview, a first date, ordering food) played completely straight in a ridiculous setting
- A character breaking an unspoken social rule and not realizing it
- Escalation — a small, reasonable thing spiraling into something ridiculous

None of these are formulas you must follow. They're starting points if you're staring at a blank page.

Strong hooks

On short-form video, you have roughly one to two seconds before someone decides whether to keep watching. Your first line, first visual, or first moment needs to do one of these:

- Create a question the viewer wants answered
- Show something visually unexpected
- Say something that immediately signals "this is going to be funny"

Bad hook (too slow):

A wide shot of a character standing in a room, then a few seconds of silence before anything happens.

Stronger hook (immediate):

The character is already mid-sentence, saying something confidently absurd, with a serious expression, as the video opens.

COMMON MISTAKE: Beginners often build up to the "funny part" slowly, like a story with a setup and a payoff. On short-form video, this usually loses the viewer before you get there. Front-load the interesting part, then fill in context after you already have their attention.

Short-form pacing

Keep the whole video short. For this style, somewhere in the range of 15–40 seconds is a reasonable target for a single video, though this isn't a strict rule — some ideas earn a longer runtime, most don't.

Practical pacing habits:

- Cut anything that doesn't move the joke forward
- If a line doesn't add a new idea, cut it
- Get to the punchline as fast as the joke allows
- End shortly after the punchline — don't over-explain it

Writing concise dialogue

Every line of dialogue should be doing one of two jobs: setting something up, or paying something off. If a line isn't doing either, it's filler.

Weak dialogue (padded):

"So, listen, I've been thinking about this for a while now, and I really think we need to talk about what happened yesterday."

Tighter dialogue (same idea, faster):

"We need to talk about yesterday."

Say less. Let the visual and the performance carry some of the weight.

Punchlines

A punchline works best when it's a natural extension of what was already set up, not a random left turn. Set up an expectation, then subvert or escalate it — don't switch to an unrelated idea at the last second, since the viewer won't have anywhere to land.

Developing your own comedic style

You will get better at this by making videos, not by reading more about it. Your early videos won't sound distinctly "yours" yet — that voice develops after you've made several. Don't wait for a perfect comedic style before you start; it comes from repetition, not planning.

Why AI shouldn't replace your creativity

AI tools are genuinely useful for brainstorming variations, polishing a line, or offering alternatives once you already have a direction. What they are not good at is originating your comedic point of view from nothing. If you ask a general-purpose AI model for "a funny script," you'll often get something technically coherent and completely generic — competent, but not actually funny in a specific way.

PRO TIP: Use AI as a brainstorming partner after you already have a rough idea, not as the source of the idea itself. "Give me five ways to escalate this premise" is a good use of AI. "Write me something funny" usually isn't.

Example premises (illustrative, not real scripts to copy)

To show what a workable premise looks like, here are a few original, hypothetical examples:

1. A character calmly explains an obviously terrible plan as if it's the most reasonable thing in the world, while another character nods along, increasingly alarmed.
2. A character treats a tiny inconvenience with the emotional intensity of a life-altering catastrophe.
3. A character confidently gives directions that are clearly, obviously wrong, while remaining completely unbothered when questioned.

These are meant only to illustrate the *shape* of a workable premise — simple, one-sentence, escalating or subverting a clear expectation. Don't copy them directly; use them to calibrate what "simple enough" looks like, then write your own.

Don't copy another creator's exact script or premise. Use other people's videos for inspiration on structure and pacing, not as a source of material.

02 — CREATE YOUR CHARACTER

Why start with the character

Your character is the anchor of the video. Before you touch environments, scripts, or shots, get your character locked down — everything else is built around them.

The tool: ChatGPT image generation

This Blueprint uses ChatGPT's image generation as the default recommendation for creating your character, because it's simple, accessible, and doesn't require specialized software. If you prefer another image-generation tool, the same principles apply — you're just using a different interface to get there.

Start from an existing reference point

A practical way to get a recognizable, interesting-looking character quickly is to start from an existing fictional character type as a reference point, then modify the design into something sufficiently your own.

People already have a mental shortcut for familiar character archetypes — a certain kind of hero, a certain kind of anime protagonist, a certain kind of video-game sidekick. Starting from something recognizable can make your character immediately readable to a viewer, even in a three-second shot.

To do this, describe an existing character type to the image model, then specify changes across categories such as:

- **Clothing** — different outfit, different colors, different era
- **Hair** — different style, length, or color
- **Colors** — shift the overall palette
- **Accessories** — add or remove props, jewelry, weapons, gear
- **Facial features** — different face shape, eyes, expression style
- **Body proportions** — taller, shorter, stockier, more exaggerated
- **Other visual details** — scars, markings, unique props

IMPORTANT — originality and commercial use: Using an existing character as a *starting reference point* for your own design process is a normal part of concept work. But if you intend to post, monetize, or sell content commercially, make sure your final design is sufficiently changed that it functions as your own original character rather than a direct copy of a copyrighted one. The more of the categories above you change, and the more you change them, the further you move from "copy" and the closer you get to "original character inspired by a type." This section is practical guidance, not legal advice — when in doubt, change more, not less.

Character Creation Template

CHARACTER CREATION TEMPLATE

Base reference type: [describe a general character archetype, e.g.
"a stylized teenage protagonist from an early-2000s adventure game"]

Changes to apply:

- Clothing: [describe]
- Hair: [describe]
- Color palette: [describe]
- Accessories: [describe]
- Facial features: [describe]
- Body proportions: [describe]
- Other distinguishing details: [describe]

Visual style: 3D rendered, low-poly, early-2000s PS2-era Japanese video-game aesthetic, baked-looking lighting, simple geometry, slightly nostalgic and imperfect, cinematic framing.

Pose/shot: [full body / close-up / side profile / etc.]

[ADD MY EXACT PROMPT HERE, if you want to swap in your own wording]

Treat this as a starting skeleton, not a rigid script. Feel free to add or remove lines depending on what you're trying to get.

Making the character visually recognizable

A design tends to read well on video if it has one or two clear, simple visual "hooks" — something a viewer can register instantly, even in motion, even at small size on a phone screen. This could be a distinctive hair silhouette, a strong color combination, or one standout accessory. Overly detailed, busy designs often read as visual noise once they're moving and small on a phone screen.

Maintaining visual consistency

Once you land on a design you like, treat it as final for that character. Don't regenerate a "new version" of the character for every scene — this is exactly what causes the character to visually drift between clips. Instead, move on to building a reference library (next chapter) using this exact design.

COMMON MISTAKE: Regenerating the character from scratch for each new scene, instead of reusing the same reference images. This is the single biggest cause of characters looking slightly different from clip to clip.

03 — CREATE THE ENVIRONMENT

The tool

Like character creation, this Blueprint uses ChatGPT image generation for environments. The same two approaches below work with any comparable image-generation tool.

Approach A — Use a visual reference

Take a screenshot — from a video you like, a game, or another suitable visual reference — and give it to the image model, asking it to create a similar environment or style while adapting it to your own scene.

This is often faster and easier than describing every visual detail from scratch, because the reference does a lot of the descriptive work for you. You're not copying the reference outright — you're using it to communicate a visual target, then asking the model to adapt it.

Approach B — Describe the environment yourself

If you don't have a reference handy, describe the environment directly. Useful things to specify:

- **Location** — what kind of place is this?
- **Camera angle** — how is it being viewed?
- **Composition** — what's in the foreground, midground, background?
- **Lighting** — time of day, light source, mood
- **Objects** — what's physically in the scene?
- **Atmosphere** — calm, tense, whimsical, eerie, etc.
- **Color palette** — dominant colors and tone
- **Perspective** — wide, tight, low, elevated
- **PS2 aesthetic** — low-poly, baked lighting, simple textures (see Chapter 11)
- **Level of detail** — sparse and stylized vs. busy and detailed
- **Character placement**, if the character needs to appear in the same image

Environment Template

ENVIRONMENT TEMPLATE

Location: [describe]

Camera angle: [describe]

Composition: [foreground / midground / background elements]

Lighting: [time of day, light source, mood]

Objects in scene: [list]

Atmosphere: [describe]

Color palette: [describe]

Perspective: [wide / tight / low / elevated]

Visual style: 3D rendered, low-poly, early-2000s PS2-era Japanese video-game aesthetic, baked-looking lighting, simple geometry, slightly nostalgic and imperfect, cinematic framing.

[ADD MY EXACT PROMPT HERE, if you want to swap in your own wording]

Design the environment around the eventual shot

Don't create an environment in a vacuum — design it with a specific shot in mind. Think about where the camera will actually be and what needs to be visible from that angle before you generate it.

You usually only need one background angle

You generally don't need to generate a separate environment image for every camera angle you plan to use. One solid environment image, generated with your main shot in mind, is often enough — your image-to-video or image-editing tool can reframe, crop, or adapt that same environment for close-ups or alternate angles later, rather than you generating an entirely new background from scratch each time.

04 — BUILD YOUR REFERENCE IMAGES

Why one image isn't enough

A single character image only shows you one pose from one angle. Later, when you're generating clips of that character walking, sitting, talking, or reacting, having a small library of reference images makes it much easier for your image-to-video tool to keep the character looking consistent.

The minimum reference set

At minimum, aim to generate:

- **Full body** — shows the complete design and proportions

- **Side profile** — useful for movement and turning shots
- **Close-up** — shows facial detail clearly
- **Different facial expressions** — at least a neutral, a surprised, and one exaggerated expression
- **Sitting** — useful for dialogue-heavy scenes
- **Leaning against something** — a common, natural-looking casual pose
- **Other useful poses** as you discover you need them (pointing, walking mid-stride, arms crossed, etc.)

You don't need to generate this entire set before you start — build it incrementally as your script and shot plan tell you what you'll actually need.

Why separate images work better than one character sheet

You do not need to force every pose into a single giant "character sheet" image. Separate images are completely fine, and are often more useful, because:

- Each individual image can be generated at higher effective detail
- You can regenerate just the pose that isn't working, without redoing the whole sheet
- Your image-to-video tool can reference the specific pose that's closest to what you need for a given shot

Character Reference Template

CHARACTER REFERENCE TEMPLATE

Character: [reference your existing character design/image]

Pose: [full body / side profile / close-up / sitting / leaning / other]

Expression: [neutral / surprised / annoyed / exaggerated / etc.]

Framing: [how tightly cropped]

Keep exact same character design, proportions, clothing, and colors as the reference. Only change pose/expression/framing as specified.

[ADD MY EXACT PROMPT HERE, if you want to swap in your own wording]

CREATOR TIP: Always explicitly tell the model to keep the same character design as your existing reference image when generating a new pose. This single instruction dramatically reduces unwanted design drift.

05 — PLAN THE SHOTS

There's no single correct system

Shot planning for this kind of content is intentionally loose. You don't need a film-school framework — you need a simple habit: look at what the moment needs, and pick a shot that serves it.

Shot type menu

Shot type	When it's useful
Close-up	Emphasizing a reaction or expression
Wide shot	Establishing a location or showing full-body action
Low angle	Making a character/moment feel more dramatic or absurdly grand
High angle	Making a character/moment feel smaller or more vulnerable
Side shot	Showing movement or a conversation between two positions
Walking shot	Following movement, building anticipation
Static/locked-off shot	Letting a performance or line land without visual distraction
Over-the-shoulder	Framing a conversation or reaction to something off-screen
Sitting shot	Calm, dialogue-focused moments
Leaning shot	Casual, relaxed character moments
Camera movement (push-in, pull-out, pan)	Building tension or emphasis toward a line or reveal

The core planning question

For every beat in your script, ask:

"What shot makes this particular moment funnier or more interesting?"

That's it. You don't need to justify the choice with film theory — trust the instinct that comes from watching videos you already like.

Take inspiration from videos you like

When you're unsure what shot to use, look at how creators you enjoy shoot a similar kind of moment. You're not copying their content — you're borrowing structural inspiration, the same way any creator studies work they admire.

Keep it simple

Don't overcomplicate this stage. A short video built from three or four well-chosen shots beats an overly ambitious plan with twelve shots you'll struggle to generate and edit together.

Hypothetical example

Say your script is: *A character calmly explains an obviously terrible plan while another character listens, increasingly alarmed.*

A simple shot plan might look like:

1. **Wide/static shot** — establish both characters in the scene
2. **Close-up on the speaker** — deliver the "terrible plan" line with a completely calm expression
3. **Close-up on the listener** — reaction shot, visibly alarmed
4. **Static two-shot** — final beat/punchline, both characters visible

Four shots. Nothing exotic. That's usually enough.

06 — GOOGLE FLOW: IMAGE TO VIDEO

What Flow does

Google Flow is an image-to-video tool: you give it a still image (your character and/or environment reference) along with a prompt describing what should happen, and it generates a short video clip.

The model matters less than the skill

Which specific Veo/Flow model is available to you may change over time — new versions get released, older ones get retired or updated. Rather than depending on one exact version, use whatever suitable model is currently available in your account. The underlying skill — prompting clearly and specifically — carries over regardless of which model version you're using.

The basic workflow

1. **Start a project** in Flow.
2. **Upload your reference image(s)** — typically your character reference and/or environment reference for the shot you're generating.
3. **Write your prompt**, describing the action, movement, camera, dialogue, and other elements for this specific clip (covered in full detail in Chapter 07).
4. **Generate** the clip.
5. **Review the result** carefully — does the character look right? Does the movement match what you described? Does the camera behave as expected?
6. **Regenerate or correct** as needed (covered in Chapter 08) if something is off.

Keep instructions general as the software evolves

Flow's interface, available models, and exact settings will likely change after this Blueprint is written. Focus on understanding *what* you're trying to communicate to the tool — a clear starting position, action, movement, camera behavior, and so on — rather than memorizing exact button locations, since those are the parts most likely to shift over time.

07 — MASTERING AI VIDEO PROMPTS

This is the most important chapter in the Blueprint. If you only deeply learn one skill from this entire product, make it this one.

The core problem

The AI cannot read your mind. If you have a specific movement, expression, or camera behavior in your head, you have to actually describe it — the model has no way to guess what you're picturing.

Weak prompt:

"Make the character walk toward the camera naturally."

This sounds reasonable, but it leaves almost everything undefined: starting position, walking speed, arm movement, facial expression, what the camera does, what happens at the end of the clip. The model will fill in all of those gaps itself, and its guesses often won't match what you imagined.

Better prompt (same idea, fully specified):

"Character starts standing still at the edge of frame, arms relaxed at their sides, neutral expression. They begin walking at a slow, casual pace directly toward the camera, taking four to five steps. Their arms swing naturally with each step. Halfway through, their expression shifts to a slight smirk. The camera remains

completely static, locked off, at chest height. Background stays unchanged. No dialogue. Ambient environment sound only, no music."

Notice what changed: every element that was previously left to chance is now explicitly defined.

The core principle

The more specific the shot, the more specific your prompt needs to be.

For a simple, low-stakes shot, a shorter prompt is fine. For a shot where a specific movement, expression, or line delivery really matters, spell it out.

The full checklist of elements to specify

When a shot matters, work through this list and decide what to specify for each:

- **Character** — who is in the shot (reference which character/image)
- **Starting position** — exactly where and how they begin
- **Action** — what they are doing
- **Movement path** — where they move, in what direction
- **Hand gestures** — what the hands/fingers are doing
- **Body movement** — torso, shoulders, head, posture
- **Facial expression** — the specific emotion or look
- **Camera** — position, and whether it's locked, handheld, tracking, or moving
- **Environment** — what should stay unchanged in the background
- **Dialogue** — the exact line(s), if any
- **Voice** — pitch, speed, tone, delivery
- **Sound** — dialogue, ambient sound, music, effects, or silence
- **Restrictions** — explicitly state what should NOT happen

The conceptual prompt structure

A useful (but flexible — not rigid) mental order to build a prompt in:

CHARACTER → STARTING POSITION → ACTION → MOVEMENT PATH → HAND
GESTURES

→ BODY MOVEMENT → FACIAL EXPRESSION → CAMERA → ENVIRONMENT

→ DIALOGUE → VOICE → SOUND → RESTRICTIONS

You won't always need every single element for every shot — a static reaction shot needs far less movement detail than a walking shot. Use this as a checklist to run through, not a form you must fill out completely every time.

Template 1 — Talking Character

Character: [reference character]
Starting position: [where they are, how they're positioned]
Action: speaking directly to camera / to another character
Facial expression: [describe]
Body movement: [minimal / gestures while speaking / etc.]
Camera: [static / slow push-in / etc.]
Environment: unchanged from reference
Dialogue: "[exact line]"
Voice: [pitch, pace, tone, delivery]
Sound: dialogue only / dialogue + ambient sound
Restrictions: [e.g., "camera does not move," "character does not walk"]

Template 2 — Walking Character

Character: [reference character]
Starting position: [location in frame, posture]
Action: walking
Movement path: [direction, distance, e.g. "walks from left edge of frame to center, then stops"]
Speed: [slow/casual, brisk, hurried]
Hand gestures: [e.g., "hands in pockets" or "arms swinging naturally"]
Body movement: [posture while walking]
Facial expression: [describe]
Camera: [static / tracking alongside / etc.]
Environment: unchanged from reference
Dialogue: [if any]
Sound: [footsteps, ambient, music, or silence]
Restrictions: [e.g., "character does not turn to face camera"]

Template 3 — Sitting Character

Character: [reference character]
Starting position: already seated, [describe posture]
Action: [e.g., "gesturing while speaking" or "remaining still"]
Facial expression: [describe]
Camera: [static, eye-level with seated character]
Environment: unchanged from reference
Dialogue: "[exact line, if any]"
Voice: [pitch, pace, tone]
Sound: [dialogue / ambient / silence]
Restrictions: [e.g., "character does not stand up"]

Template 4 — Character Leaning

Character: [reference character]

Starting position: leaning against [object/surface from environment],
[describe arm/leg position]

Action: [e.g., "remains in leaning pose, arms crossed" or "gestures casually while speaking"]

Facial expression: [describe]

Camera: [static / slow drift]

Environment: unchanged from reference

Dialogue: [if any]

Sound: [ambient / dialogue / silence]

Restrictions: [e.g., "character does not change position"]

Template 5 — Character Approaching Camera

Character: [reference character]

Starting position: standing still, [distance from camera]

Action: walking directly toward the camera

Movement path: straight line, ending at [describe end position/distance]

Speed: [slow and deliberate / brisk]

Hand gestures: [describe]

Body movement: [describe]

Facial expression: [describe, and whether it changes mid-walk]

Camera: locked off, static, [height]

Environment: unchanged from reference

Dialogue: [if any]

Sound: [ambient / dialogue / silence]

Restrictions: character does not stop before reaching described end position; camera does not move

Template 6 — Close-Up

Character: [reference character]

Framing: tight close-up on face/upper body

Starting expression: [describe]

Action: [e.g., "expression shifts from neutral to surprised" or "remains still"]

Camera: static, locked off

Environment: unchanged (mostly out of frame)

Dialogue: "[exact line, if any]"

Voice: [pitch, pace, tone, delivery]

Sound: [dialogue / ambient / silence]

Restrictions: [e.g., "character does not turn head away from camera"]

Template 7 — Two-Character Interaction

Characters: [reference character A] and [reference character B]

Starting position: [describe positions relative to each other]

Action: [describe what each character is doing]

Character A dialogue: "[exact line]"

Character A expression/voice: [describe]

Character B reaction: [describe expression/body language]

Camera: [static two-shot / over-the-shoulder / etc.]

Environment: unchanged from reference

Sound: [dialogue / ambient / silence]

Restrictions: [e.g., "neither character changes position during the clip"]

Template 8 — Movement Correction

Keep: [describe exactly what worked and should stay the same — character design, environment, camera, overall pacing, etc.]

Fix: [describe precisely what was wrong and what should happen instead]

Do not change: [explicitly list anything that should remain untouched]

PRO TIP: These templates are starting skeletons, not sacred formulas. Delete lines that don't apply to your shot. Add lines for anything specific to your scene that isn't covered above.

COMMON MISTAKE: Writing an enormous, over-stuffed prompt "just in case," describing details that don't actually matter for that shot. This can dilute the parts that do matter. Be specific about what matters, and don't pad the rest.

08 — REGENERATING BAD GENERATIONS

The iteration loop

Your first generation will often not be exactly right. This is completely normal and not a sign you're doing something wrong. The process is:

GENERATE → INSPECT → IDENTIFY THE PROBLEM → CORRECT THE PROMPT → REGENERATE

Identify the specific mistake

The most important habit here is precision. Don't react to a bad generation by rewriting the entire prompt from scratch — figure out exactly what's wrong first.

Common specific problems:

- Wrong hand movement
- Character turns too quickly
- Character walks in the wrong direction
- Voice too high or too low
- Voice too slow or too fast
- Camera moves unexpectedly
- Character stops too early
- Wrong facial expression
- Character changes appearance (design drift)
- Background changes unexpectedly
- Movement looks unnatural/floaty

Correct only what's wrong

Once you know the specific issue, write a targeted correction rather than a full rewrite:

Example — wrong movement direction:

"Keep the same character, environment, and camera setup. Fix: the character walked to the right instead of straight toward the camera. Change the movement path so the character walks in a straight line directly toward the camera. Do not change anything else about the shot."

Example — voice issue:

"Keep the same voice identity and delivery style. Fix: the pitch was too high. Lower the pitch slightly and keep everything else about the voice the same."

Example — design drift:

"Keep the exact same character design, clothing, and colors as the original reference image. Fix: the character's outfit color changed. Regenerate using the original reference image as the visual source of truth."

QUICK FIX: If a generation is close but one small thing is off, always try a targeted correction prompt before regenerating from scratch. It's faster, cheaper, and more likely to preserve the parts that were already working.

09 — VOICE & PERFORMANCE

Two options: native voice vs. dedicated tools

Option 1 — Flow's native voice generation. This is the simplest path: your image-to-video tool generates dialogue and voice as part of the same clip. It's convenient, but voice characteristics can vary somewhat between separate generations, since each clip is generated independently.

Option 2 — A dedicated voice-generation tool, such as ElevenLabs. For more control and more consistency across multiple clips, consider generating your voice separately using a dedicated tool, then combining that audio with your video clips during editing. Dedicated voice tools are generally built specifically for consistent voice identity across many generations, which can be harder to guarantee from a general-purpose video tool's built-in voice.

Neither approach is "wrong" — native voice is faster to get started with; a dedicated tool gives you more control once consistency matters more to you. Specific features, models, and pricing for any voice tool can change over time, so treat this as a general direction rather than a permanent recommendation of one exact product feature.

What you can control in a voice

Whichever approach you use, you can describe and adjust:

- **Voice identity** — the underlying character/quality of the voice
- **Pitch** — higher or lower
- **Speed** — faster or slower delivery
- **Enunciation** — how clearly words are pronounced
- **Delivery style** — deadpan, energetic, sarcastic, sincere, etc.
- **Tone** — warm, cold, flat, dramatic
- **Emotional intensity** — subdued vs. exaggerated
- **Consistency** — keeping the same identity across multiple clips

Combining a separate voice with your video

If you generate voice separately, the general workflow is: generate the video clip (with or without placeholder/no dialogue), generate your voice performance separately, then bring both into CapCut and align them during editing (covered in Chapter 10).

Voice Prompt Template

Voice identity: [describe — e.g., "young adult male, slightly nasal"]

Pitch: [higher / lower / neutral]

Speed: [slower / faster / neutral]

Delivery: [deadpan / energetic / sarcastic / sincere / etc.]

Tone: [describe]

Emotional intensity: [subdued / exaggerated / neutral]

Line: "[exact dialogue]"

[ADD MY EXACT VOICE PROMPT HERE, if you want to swap in your own wording]

Correcting a voice

The same specificity principle from Chapter 07 applies directly to voice.

Weak correction:

"Make the voice better."

Specific correction:

"Keep the same voice identity, but lower the pitch slightly, speak faster, use shorter pauses, and maintain a calm deadpan delivery."

Useful properties to reference when correcting a voice:

- Higher/lower pitch
- Faster/slower speech
- Calmer/more aggressive
- Deadpan/more emotional
- Relaxed/more intense
- Clearer/rougher
- Shorter/longer pauses
- Pronunciation of a specific word
- More/less enunciation
- Higher/lower energy

Voice references

Some tools support providing a reference voice recording to guide generation. This isn't a required part of the core workflow taught here — it's an optional feature worth exploring later, depending on which tool and subscription tier you're using, once you're comfortable with the simpler approach of describing the voice you want directly.

10 — BUILDING MULTIPLE CLIPS

Why short clips beat one long generation

Don't try to generate one long, continuous AI video. Instead, build your final video out of several short clips. A useful practical target is **around 8–10 seconds or less per clip when appropriate** — though treat this as a practical guideline, not a fixed technical limit, since exact tool capabilities change over time.

Reasons short clips work better:

- **Easier to regenerate** — redoing 8 seconds is faster than redoing 40
- **Easier to fix** — a small, isolated mistake is easier to identify and correct
- **Easier to maintain consistency** — less time for the character or environment to visually drift within a single generation
- **Less wasted credits/generation time** — a bad short clip costs less than a bad long one
- **Easier to edit** — short clips give you more flexibility and control when assembling in CapCut
- **One mistake doesn't ruin an entire scene** — you only need to replace the specific clip that went wrong, not the whole thing

Breaking your script into clips

Take your shot plan from Chapter 05 and treat each shot as its own separate generation. A four-shot script becomes four separate short clips, generated and reviewed individually.

Maintaining consistency across clips

Use the reference image library you built in Chapter 04 consistently across every clip featuring that character. Don't regenerate a "new" character reference partway through — reuse the same original images as your source of truth for the character's appearance throughout the entire video.

Example: one script becoming several clips

Using the shot plan from Chapter 05's example (a character explaining a terrible plan):

1. **Clip 1** — wide/static establishing shot of both characters (no dialogue, ~3–4 seconds)
2. **Clip 2** — close-up on the speaker delivering the "plan" line, calm expression (~6–8 seconds)
3. **Clip 3** — close-up reaction shot on the listener, visibly alarmed (~4–5 seconds)
4. **Clip 4** — static two-shot for the final punchline beat (~5–6 seconds)

Four separate, short generations. Each one gets reviewed and, if necessary, corrected on its own before moving to the next.

Avoiding continuity problems

Keep a simple mental (or written) note of what each character is wearing, holding, and doing at the end of one clip versus the start of the next, so clips cut together believably. If clip 2 ends with a character's arms crossed, clip 3 shouldn't open with their arms at their sides unless that change is intentional.

11 — CAPCUT EDITING

Keep it simple

The goal here is a finished, watchable short-form video — not an advanced filmmaking exercise. CapCut is recommended for its accessibility, but any editor you're comfortable with will work the same way conceptually.

The editing workflow

1. **Import your clips** into a new CapCut project.
2. **Put them in the correct order** on the timeline, matching your shot plan.
3. **Trim unnecessary parts** — cut any dead air, false starts, or extra frames at the beginning/end of each clip.
4. **Fix pacing** — watch the sequence back and tighten anything that drags.
5. **Add music**, if appropriate for the video's tone — keep it subtle enough that it doesn't compete with dialogue.
6. **Add captions** — especially important since much of short-form viewing happens with sound off.
7. **Add sound effects**, where they'd genuinely add impact (a footstep, a comedic sting, etc.) — don't overdo it.
8. **Add transitions**, only where they're actually needed — most cuts in this style of video work fine as simple hard cuts.
9. **Add simple zoom/camera effects**, if a moment would benefit from a subtle push-in for emphasis.
10. **Adjust final pacing** one more time after everything is assembled.
11. **Export** your finished video.

CREATOR TIP: Watch your rough cut all the way through at least once before adding captions, music, or effects. Pacing problems are much easier to spot on a clean cut than once it's fully decorated.

COMMON MISTAKE: Adding too many transitions and effects to "make it feel more produced." For this style of short-form video, simple, clean cuts usually serve the comedy better than heavy editing flourishes.

12 — THE PS2 LOOK

It's simpler than it looks

The PS2-era aesthetic can feel like a specific, mysterious style to pin down, but once you understand what you're actually recreating, it becomes much easier to describe and generate consistently.

The visual target

- **Low-poly 3D geometry** — simplified shapes, not hyper-detailed models
- **Lower-resolution textures** — slightly soft, less crisp than modern 3D
- **Simple geometry** — fewer fine details on objects and environments
- **Baked-looking lighting** — lighting that looks pre-computed and static rather than fully dynamic/realistic
- **Early-2000s Japanese video-game aesthetic** — the general visual language of that console generation
- **Slight visual imperfections** — a touch of roughness rather than a hyper-polished modern render
- **Nostalgic atmosphere** — a soft, slightly hazy or warm quality
- **Cinematic framing** — composed like a film shot, not a flat game-menu screenshot
- **Game-like character/environment design** — stylized rather than photorealistic

The easiest way to achieve it consistently

The same reference-image method from earlier chapters applies directly here:

Find a visual reference you like → give it to the image model → ask for a similar PS2-era aesthetic → adjust it for your specific scene.

Don't overthink it

It's common to spend far too much time experimenting with style keywords, chasing a "perfect" look. Once you understand the general visual language above, focus more of your energy on consistency (using the same references across your project) than on endless style experimentation.

PS2-Style Prompt Add-On

Add this to any character or environment prompt from earlier chapters:

Visual style: 3D rendered, low-poly geometry, lower-resolution textures, baked-looking lighting, simple geometry, early-2000s Japanese video-game aesthetic, slightly nostalgic and imperfect,

cinematic composition and framing.

13 — COMPLETE VIDEO CASE STUDY

This is an illustrative, hypothetical example built to demonstrate how every chapter connects — it is not a claim about any specific real video.

IDEA

A character calmly explains an obviously terrible plan to a friend, who becomes increasingly alarmed while the first character remains completely unbothered.

CHARACTER

Using the Character Creation Template from Chapter 02: a stylized young adult character, based on a generic "confident anime protagonist" archetype, modified with a distinct color palette, a simplified hairstyle, and one clear accessory (e.g., a jacket with an unusual color combination) to make the design original and instantly readable.

REFERENCE IMAGES

Following Chapter 04: a full-body reference, a side profile, a neutral close-up, a calm/confident expression close-up, and a sitting pose — enough coverage for the shots planned below.

ENVIRONMENT

Following Chapter 03: a simple interior location (e.g., a plain living room), generated using Approach B (text description), specifying a static wide angle, warm late-afternoon lighting, minimal background objects, and the PS2-style visual add-on from Chapter 12.

SCRIPT

Speaker: "Okay, so the plan is: we tell them it already happened." Listener: [no line — reaction only] Speaker: "It's foolproof."

SHOT PLAN

Following Chapter 05:

1. Wide static shot — both characters in the room
2. Close-up on speaker — delivers the "plan" line, calm expression

3. Close-up on listener — silent, alarmed reaction
4. Close-up on speaker again — delivers "It's foolproof," completely unbothered

FLOW PROMPT (Clip 2 example, using Template 6 — Close-Up)

Character: [speaker reference image]

Framing: tight close-up on face/upper body

Starting expression: neutral, relaxed

Action: speaking directly to another character just off-frame

Camera: static, locked off

Environment: unchanged from reference (soft living room background, out of focus)

Dialogue: "Okay, so the plan is: we tell them it already happened."

Voice: calm, medium pitch, unhurried pace, completely sincere delivery

Sound: dialogue only, quiet ambient room tone underneath

Restrictions: character does not smile or break composure; camera does not move

GENERATED CLIP

The resulting ~6-second clip shows the speaker delivering the line with a locked-off camera and calm delivery, matching the prompt.

VOICE

In this example, the native voice generation from the video tool was used directly, since consistency across only two speaking clips (from the same character) was manageable without a separate dedicated voice tool.

CAPCUT

All four clips imported and trimmed to remove dead frames at the start/end of each. Ordered per the shot plan. Captions added to both dialogue lines. A single subtle comedic sound sting added right after "It's foolproof." No music, since the dialogue and reactions carry the pacing on their own. Simple hard cuts between all four shots — no transitions needed.

FINAL RESULT

A four-shot, roughly 20-second video: establish → setup line → silent reaction → punchline. Every chapter of this Blueprint touched at least one part of this short sequence — this is the same basic shape you'll repeat for your own videos.

14 — PROMPT LIBRARY

A condensed, copy-ready collection of templates. Replace anything in [brackets].

Character creation

Base reference type: [archetype]

Changes: clothing [], hair [], colors [], accessories [],
facial features [], proportions [], other []

Visual style: 3D, low-poly, PS2-era aesthetic, baked lighting,
cinematic framing

Character modifications

Keep: exact same character design, proportions, and colors as reference

Change: [specific element only]

Character reference images

Character: [reference existing design]

Pose: [full body / profile / close-up / sitting / leaning / other]

Expression: [describe]

Keep exact same character design as reference. Only change pose/expression.

Environment

Location: []

Camera angle: []

Lighting: []

Objects: []

Atmosphere: []

Palette: []

Visual style: 3D, low-poly, PS2-era aesthetic, baked lighting,
cinematic framing

PS2 style add-on

Visual style: 3D rendered, low-poly geometry, lower-resolution
textures, baked-looking lighting, simple geometry, early-2000s
Japanese video-game aesthetic, cinematic composition.

Talking

Character: [reference]

Action: speaking [to camera / to another character]

Expression: []

Camera: [static / push-in]

Dialogue: "[line]"

Voice: [pitch/pace/tone/delivery]

Restrictions: []

Walking

Character: [reference]

Movement path: [from/to]

Speed: []

Hands: []

Expression: []

Camera: [static / tracking]

Restrictions: []

Sitting

Character: [reference]

Posture: []

Action: []

Expression: []

Camera: [static, eye-level]

Restrictions: []

Leaning

Character: [reference]

Leaning against: [object/surface]

Posture: []

Expression: []

Restrictions: []

Camera movement

Camera type: [locked-off / push-in / pull-out / pan / tracking]

Speed of movement: []

Start position: []

End position: []

Restrictions: [e.g., "no handheld shake"]

Close-up

Character: [reference]

Framing: tight close-up

Expression: [start] → [end, if it changes]

Camera: static

Restrictions: []

Voice / dialogue

Voice identity: []

Pitch: []

Speed: []

Delivery: []

Line: "[dialogue]"

Movement correction

Keep: [what worked]

Fix: [specific problem]

Do not change: []

Consistency correction

Keep exact same character design, clothing, and colors as original reference image. Fix: [describe drift]. Regenerate using the original reference as the visual source of truth.

Background correction

Keep exact same environment, lighting, and layout as original reference. Fix: [describe unwanted background change].

Facial-expression correction

Keep: [everything else about the shot]

Fix: expression should be [describe correct expression] instead of [describe what it currently is]

15 — TROUBLESHOOTING

Problem	Why it happens	What to change	Example correction prompt
Character changes appearance between generations	Each generation is created somewhat independently; without a strong reference, small design details can drift	Always reference your original character image explicitly and instruct the model to keep the exact design	"Keep exact same character design, clothing, and colors as the original reference image. Regenerate using that reference as the visual source of truth."
Hands/fingers look wrong	Hands are a commonly difficult area for AI generation in general	Simplify the hand action requested, or specify a pose that keeps hands less prominent/detailed (e.g., in pockets, at sides)	"Keep everything the same. Change: hands remain relaxed at the character's sides instead of gesturing."
Movement is unnatural	Prompt likely under-specified the movement details	Add specific detail about speed, path, and body movement	"Keep the same shot. Change: movement should be a slow, deliberate walk with natural arm swing, not a rushed or jerky motion."
Character walks incorrectly (wrong direction/speed)	Movement path wasn't explicitly defined	Specify exact direction, path, and speed	"Keep everything the same. Change: character should walk in a straight line directly toward the camera, at a slow, casual pace."
Camera moves when it shouldn't	"Static" or "locked-off" wasn't explicitly stated	Add an explicit camera restriction	"Keep the same shot. Change: camera must remain completely static and locked off — no movement at all."

Voice changes between clips	Voice generated independently each time without a consistency anchor	Use a dedicated voice tool for multi-clip consistency, or explicitly reference the previous voice's qualities	"Keep the same voice identity, pitch, and delivery style as the previous clip featuring this character."
Words are mispronounced or cut off	Dialogue may be too long, complex, or ambiguous for the model to render clearly	Shorten the line, simplify wording, or specify pacing	"Keep everything the same. Change: slow down the delivery slightly and clearly enunciate the word '[word]'."
Environment changes unexpectedly	Background wasn't explicitly locked as "unchanged"	Add an explicit environment restriction	"Keep the same environment. Change: background must remain exactly as shown in the reference image — no new objects or layout changes."
Facial expression doesn't match the scene	Expression wasn't explicitly specified, or was under-described	Describe the specific expression more precisely	"Keep everything the same. Change: expression should be a subtle, unimpressed look rather than neutral."
Video feels slow (pacing issue)	Clips run too long, or contain dead time at the start/end	Trim more aggressively in CapCut; regenerate the clip shorter if the AI itself paced it slowly	Re-generate with: "Keep everything the same. Change: shorten the pause before the character begins speaking."
AI ignores part of your instructions	The prompt may be too long, too vague, or contains conflicting instructions	Simplify the prompt, remove conflicting details, and clearly prioritize the one or two things that matter most for this shot	Break the correction into a single, focused instruction rather than several at once: "Keep everything the same. Change only: [one specific thing]."

PRO TIP: When multiple things are wrong with a generation, fix the single most important one first and regenerate, rather than trying to correct everything at once. Compounding corrections in one prompt makes it harder for the model to isolate what you actually want changed.

16 — FINAL CHECKLIST

IDEA

- ☐ Strong hook in the first 1–2 seconds
- ☐ Original premise (not copied from another creator)
- ☐ Short, tight script
- ☐ Clear punchline

CHARACTER

- ☐ Character design finalized and sufficiently original
- ☐ Full-body reference created
- ☐ Side-profile reference created
- ☐ Close-up reference created
- ☐ Useful pose references created (sitting, leaning, expressions)

ENVIRONMENT

- ☐ Background created for your main shot
- ☐ Camera angle matches planned shot
- ☐ Lighting matches intended mood
- ☐ PS2 aesthetic applied consistently

VIDEO

- ☐ Shots planned for each script beat
- ☐ Prompts written with full detail for shots that matter
- ☐ Clips kept short (~8–10 seconds or less, where appropriate)
- ☐ Movement reviewed on each clip
- ☐ Camera behavior reviewed on each clip
- ☐ Character consistency checked across all clips

VOICE

- ☐ Voice approach chosen (native or dedicated tool)
- ☐ Delivery matches the tone of each line
- ☐ Enunciation is clear

- [] Voice consistency checked across clips featuring the same character

EDITING

- [] All clips assembled in correct order
 - [] Pacing tightened, dead time trimmed
 - [] Captions added
 - [] Music added, if appropriate
 - [] Sound effects added, if appropriate
 - [] Final export reviewed before posting
-

BONUS — MAKE YOUR FIRST VIDEO TONIGHT

If you want the fastest possible path to your first finished video, here it is, stripped down:

1. **Think of an idea** — one sentence, simple, with a clear hook.
2. **Create a character** — one solid design you're happy with.
3. **Create 3–6 character reference images** — enough to cover your planned shots.
4. **Create one environment** — designed around your main shot.
5. **Write a short script** — a few lines, tight, no filler.
6. **Break it into shots** — three or four is plenty for a first attempt.
7. **Generate each shot in Flow** — one at a time, reviewing as you go.
8. **Fix mistakes** — use targeted corrections, not full rewrites.
9. **Generate/prepare the voice** — native voice is fine to start.
10. **Assemble everything in CapCut** — order, trim, captions, minimal effects.
11. **Export.**

That's the entire Blueprint, compressed into one checklist. Everything else in this document exists to help you do each of these eleven steps well.

BONUS — CREATOR MINDSET

Don't chase perfection.

Your first video probably won't be great. That's fine — it's not supposed to be. The goal of your first video isn't to be your best work; it's to actually understand the workflow by running through it once, start to finish.

Every step in this Blueprint gets faster and more intuitive after you've done it a few times. The character-creation process that feels slow the first time becomes quick by your third character.

The prompting that feels awkward at first becomes second nature once you've corrected a dozen bad generations.

Make something. Post it. Make the next one better. That's the actual process — not a secret shortcut, just repetition applied to a workflow that works.

A NOTE ON COPYRIGHT AND COMMERCIAL USE

(This is practical guidance, not legal advice.)

Using an existing fictional character as a starting reference point can be a genuinely useful part of the ideation and learning process. However, keep the following in mind, especially if you plan to post, sell, or otherwise commercially use your videos:

- Directly reproducing a recognizable copyrighted character, largely unchanged, can create legal and intellectual-property risks.
- For commercial or public-facing work, modify your references enough that the result functions as your own original character, not a copy of an existing one.
- Respect trademarks and likeness rights — this applies to fictional characters as well as real people.
- Respect the terms of service of whichever AI tools you use, since these can include their own rules about generated content and commercial use.

If you're ever unsure whether a specific design or use case is a problem, that uncertainty itself is a signal to change more of the design or seek specific legal advice for your situation — this section is meant to flag the general risk, not resolve every individual case.
